

---

# **Pytorch Lightning**

## **Model Monitoring**

## **Hyper-parameter optimization**

---

# Pytorch Lightning



# Clean and Modular Code

PyTorch Lightning enforces a standard, modular structure:

- `forward()` → defines the model's forward pass.
- `training_step()` → defines one training iteration.
- `validation_step()` → defines validation logic.
- `configure_optimizers()` → sets up optimizers/schedulers.

This separation keeps research logic (model) and engineering logic (training loop) cleanly decoupled.

<https://lightning.ai/lightning-ai/studios/image-segmentation-with-pytorch-lightning>

---

## **Lightning handles the repetitive “boilerplate” code**

- Training and validation loops
- Epoch and batch management
- Logging and checkpointing
- GPU/TPU/multi-node setup



# Lightning Module and Callbacks

[https://lightning.ai/docs/pytorch/LTS/common/lightning\\_module.html](https://lightning.ai/docs/pytorch/LTS/common/lightning_module.html)

<https://lightning.ai/docs/pytorch/stable/extensions/callbacks.html>

[https://lightning.ai/docs/pytorch/stable/api\\_references.html#callbacks](https://lightning.ai/docs/pytorch/stable/api_references.html#callbacks)

---

## Multi-GPU and TPU Training

```
trainer = pl.Trainer(accelerator="gpu", devices=4)
```

---

## Built-in Logging and Checkpointing

More on that later



## Reproducibility and Debugging Tools

Provides utilities for deterministic training and seed control.

`Trainer(fast_dev_run=True)` allows you to **debug your pipeline quickly** using just one batch.

PyTorch Lightning = “PyTorch, but cleaner, faster, and easier to scale.”

---

# Model monitoring

---



**When do we monitor DL models? (In what stages of life-cycle?)**

---

## **When do we monitor DL models? (In what stages of life-cycle?)**

- Training
- Inference

---

# Model training

---



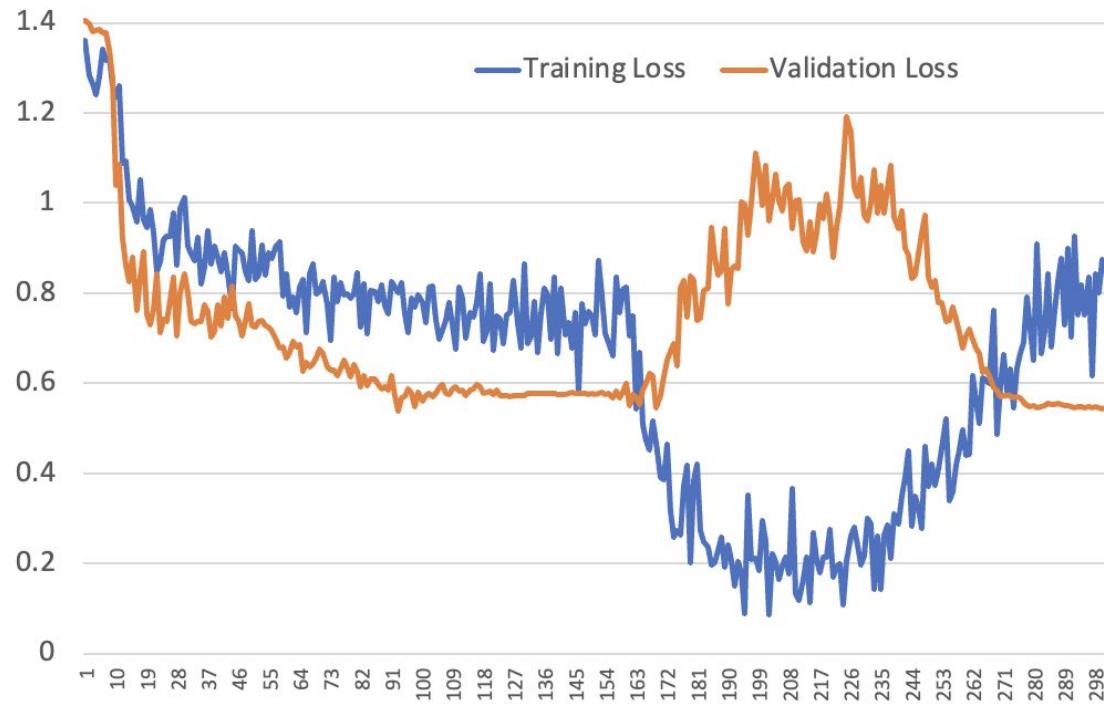
## **What and why can we monitor during DL training?**

-

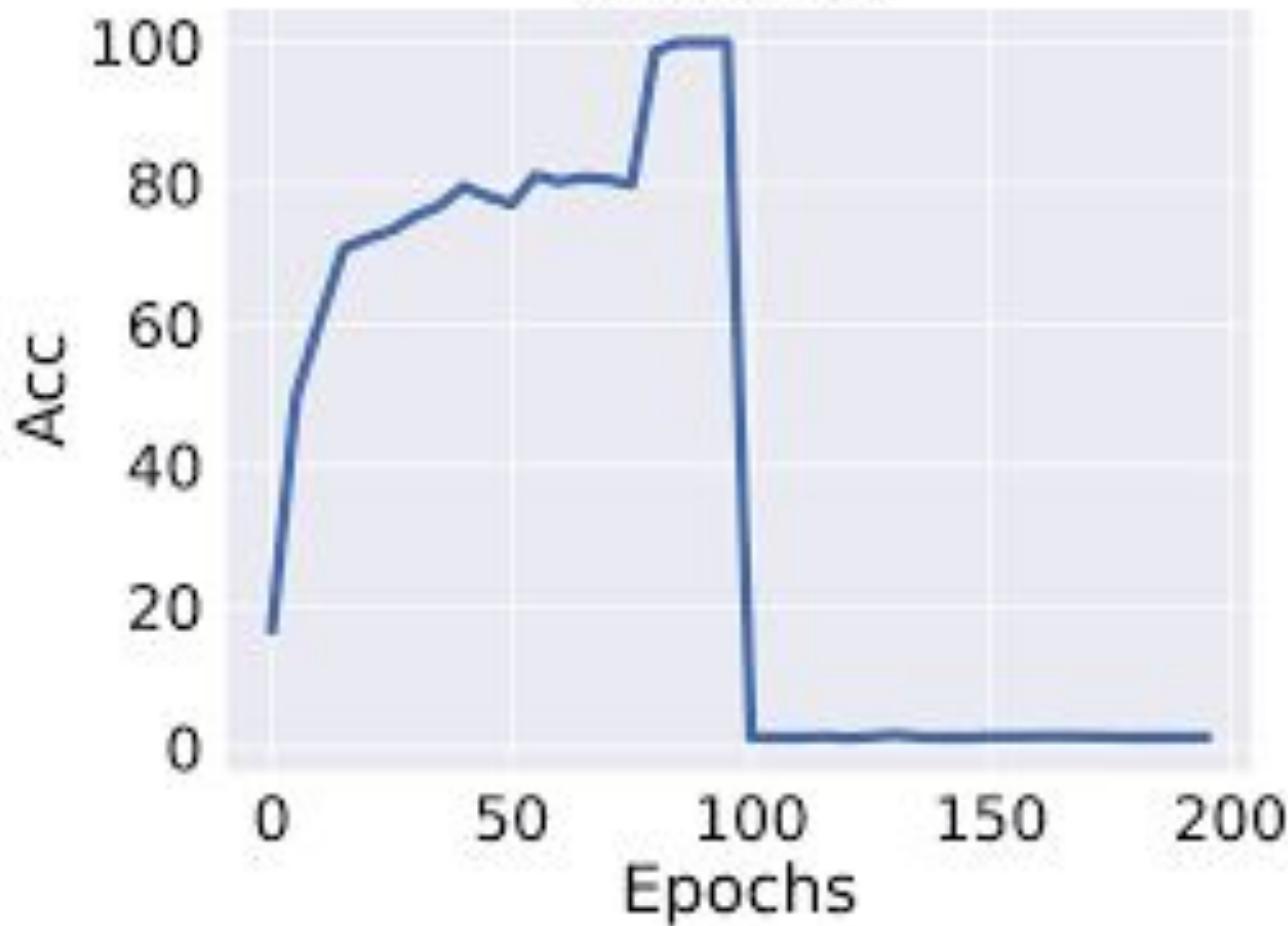
---

## What and why can we monitor during training?

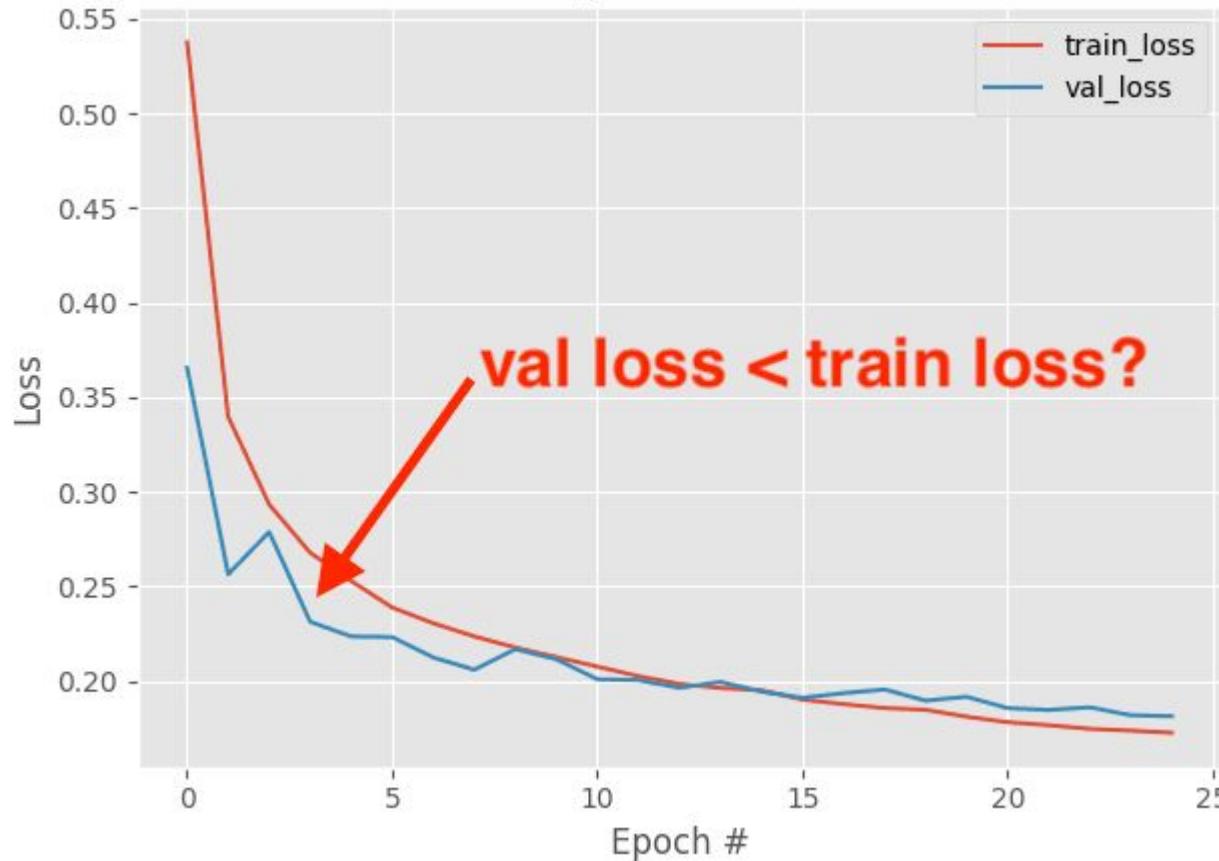
- Training and validation **loss** -> Is model learning?
- Training and validation **metrics**
- Number of **batches** (steps) and consequently **epochs**
- Log **hyper-parameters**:
  - Monitor the one that change: **Learning rate** / Learning rates
- **Hardware** usage
- I encourage you to **visualize sample** predictions!!
- At the end, all the same but for test set.



CIFAR100



### Training/Validation Loss





# TensorBoard



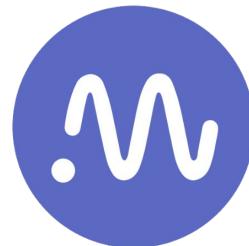
---

## Examples of tools used for model monitoring



## Weights & Biases

mlflow™



## neptune.ai



---

## Using model training monitoring tools

<https://docs.wandb.ai/models/tutorials/pytorch>

<https://docs.wandb.ai/models/integrations/lightning>

---

## Live example of results

---

# Model inference

---

## **Why do we have to monitor model after deployment?**

What should we monitor?

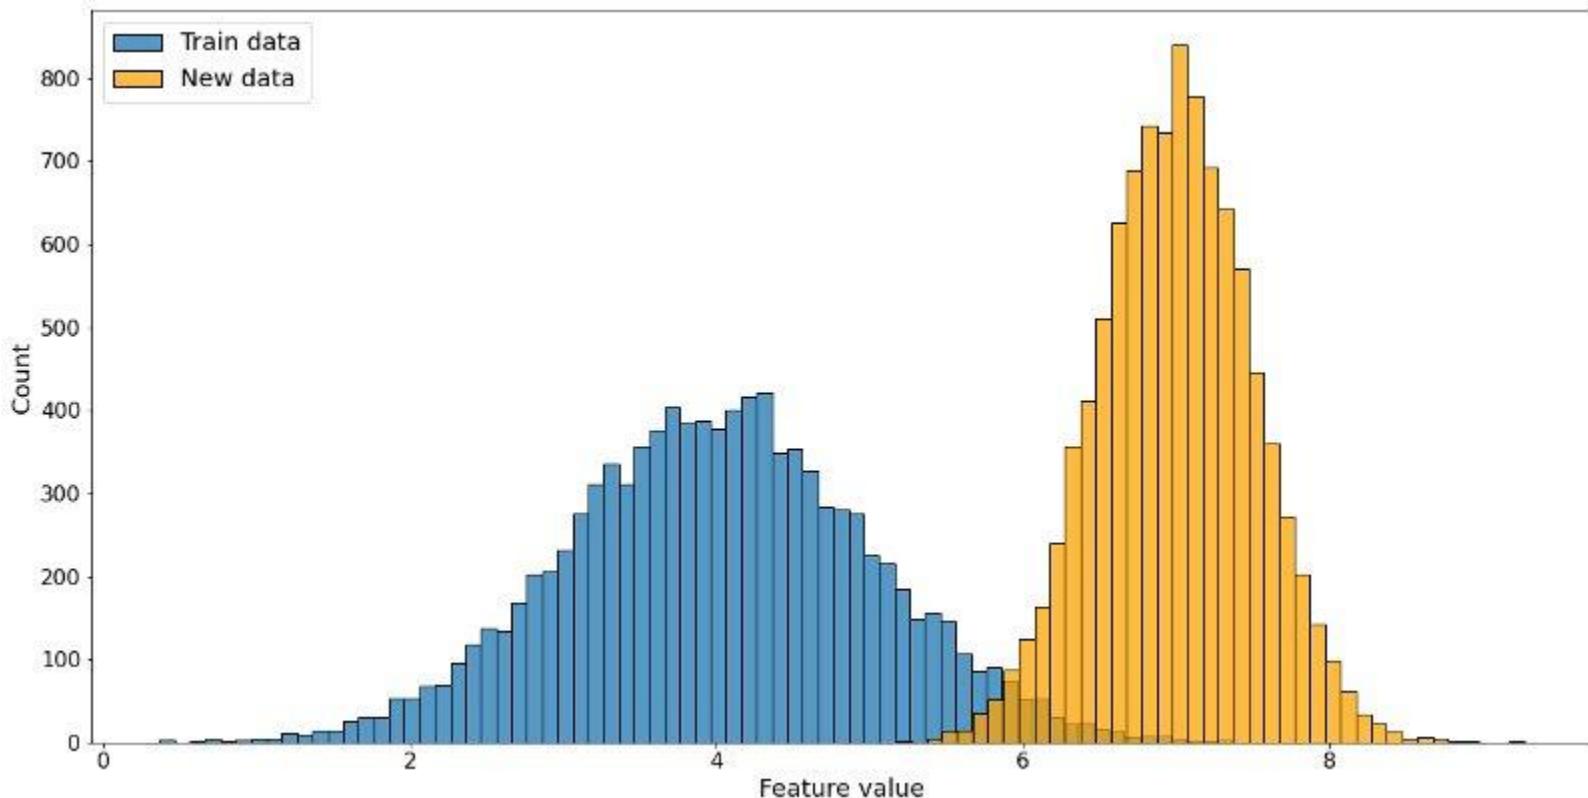
---

# Why do we have to monitor model after deployment?

What should we monitor?

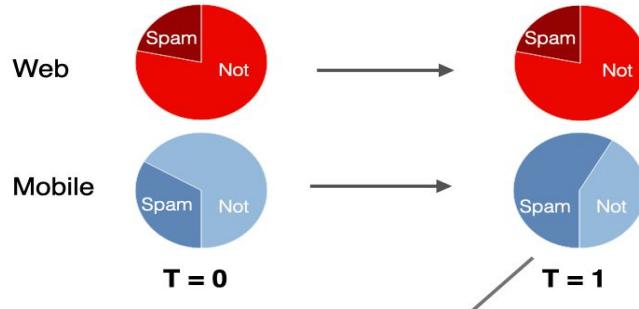
- Distribution of **input**
- Distribution of **outputs**

# Data drift

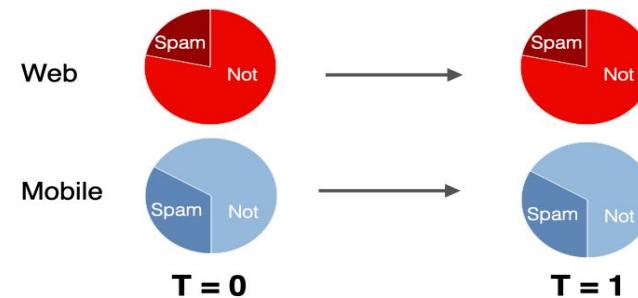


# Concept drift

## Concept drift



## Data drift



“Device type” feature distribution over time



“Device type” feature distribution over time



<https://docs.aws.amazon.com/sagemaker/latest/dg/clarify-model-monitor-feature-attribution-drift.html>



## Feature Attribution Drift

College Admission Hypothetical Scenario

Feature	Attribution in training data	Attribution in live data
SAT score	0.70	0.10
GPA	0.50	0.20
Class rank	0.05	0.70

# Feature Attribution Metrics

**SHAP (SHapley Additive exPlanations)**

**LIME (Local Interpretable Model-agnostic Explanations)**

**Integrated Gradients** (for deep models)

## Option A — Rank correlation

Compute Spearman's rank correlation coefficient between the training and live feature importance rankings.

python

 Copy code

```
import pandas as pd
from scipy.stats import spearmanr

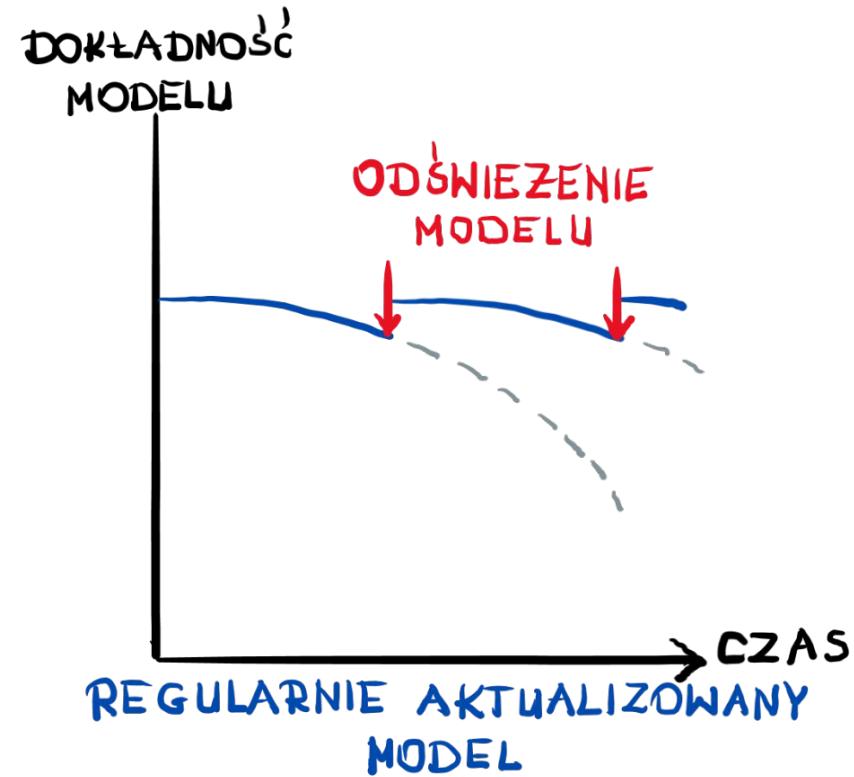
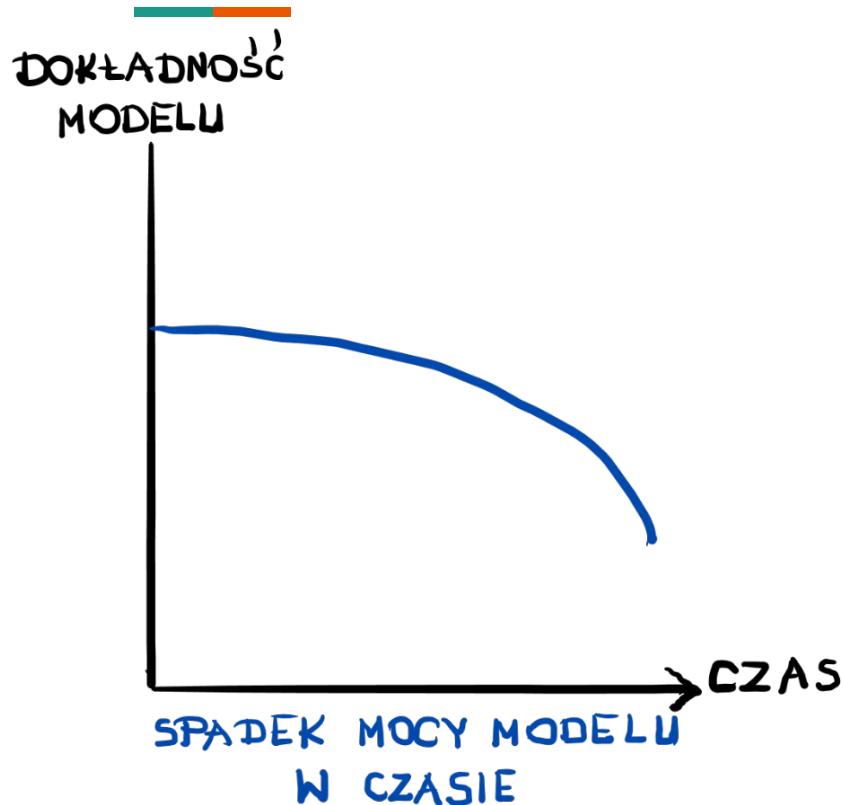
train = pd.Series({'SAT': 0.70, 'GPA': 0.50, 'Rank': 0.05})
live = pd.Series({'SAT': 0.10, 'GPA': 0.20, 'Rank': 0.70})

corr, _ = spearmanr(train.rank(), live.rank())
print(f"Spearman correlation: {corr:.2f}")
```

---

## **What to do if the model is failing?**

# Model drift = Concept Drift + Data Drift



## ml.vish\_mm\_db\_v2.ml\_28386\_classification Monitoring

+ Add tag

Share

Schedule

Refresh

⋮

**Regression**

- MSE
- RMSE
- MAE
- MAPE
- r2\_score

**Classification**

- accuracy\_score
- precision
- recall
- f1\_score
- confusion\_matrix

**Summary statistics**

- count
- distinct\_count
- quantiles
- min\_size
- max\_size
- avg\_size
- frequent\_items

**Categorical**

- min\_len
- max\_len
- avg\_len

**Numeric**

- avg
- min
- median
- max
- distinct\_count
- stddev

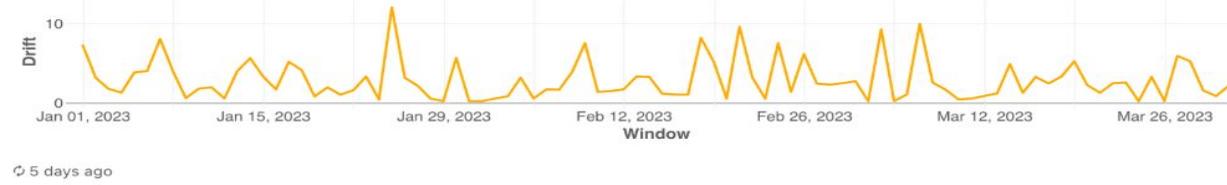
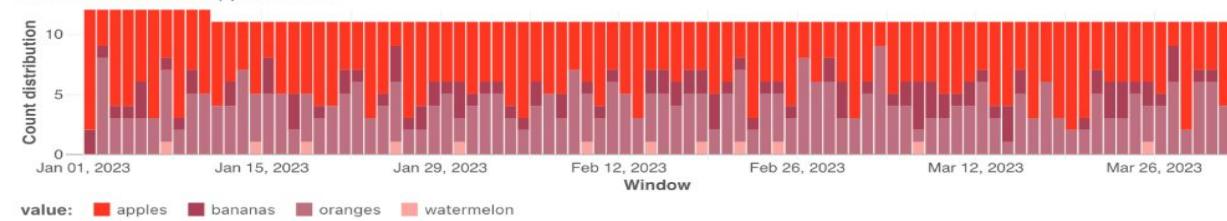
**Data integrity**

- percent\_nulls
- percent\_zeros
- percent\_nan

**Data integrity, predictions****Drift, predictions**

Compare to Metric

Baseline chi\_squared\_test.statistic

**Distribution over time, predictions**

# databricks

<https://github.com/aws-samples/amazon-sagemaker-drift-detection>

<https://aws.amazon.com/blogs/machine-learning/bring-your-own-container-to-project-model-accuracy-drift-with-amazon-sagemaker-model-monitor/>

---

# Hyper-parameters

---

## Hyper-parameter optimization

- What kind of hyper-parameters there are?
- And why do we want to know them?
- And why do we want to observe them in subsequent training runs?

---

# Hyper-parameters

- Learning rate
- Optimizers and their parameters, schedulers and their parameters
- Transformations and their parameters
- Architectures type, activation functions, number of layers, freezing, backbone types, head types
- Loss functions and their parameters
- Training Procedure: Number of epochs, gradient accumulation
- Presence and type of regularization
- Let's look at my Wandb for more examples

---

## Hyper-parameter optimization approaches

- Grid search (everything with everything)
- Random search
- Bayesian optimization
- Tree-structured Parzen Estimator
- Genetic: For Optuna: NSGA-II

---

## Tools for hyper-parameter optimization



[https://github.com/optuna/optuna-examples/blob/main/pytorch/pytorch\\_lightning\\_simple.py](https://github.com/optuna/optuna-examples/blob/main/pytorch/pytorch_lightning_simple.py)

---



**Thank you for attention**





- Why use it?

- “*PyTorch Lightning is just organized PyTorch - Lightning disentangles PyTorch code to decouple the science from the engineering.*”
- <https://lightning.ai/lightning-ai/studios/image-segmentation-with-pytorch-lightning>
- [https://lightning.ai/docs/pytorch/stable/api\\_references.html#callbacks](https://lightning.ai/docs/pytorch/stable/api_references.html#callbacks)
- [https://lightning.ai/docs/pytorch/stable/notebooks/course\\_UvA-DL/07-deep-energy-based-generative-models.html#Sampling-buffer](https://lightning.ai/docs/pytorch/stable/notebooks/course_UvA-DL/07-deep-energy-based-generative-models.html#Sampling-buffer)